Communication

Date	21 October 2026
Team ID	TDS-AI-2026
Project Name	Personalized Travel Itinerary Planner
Maximum Marks	3 Marks

Project Title

Personalized Travel Itinerary Planner

Communication During Demonstration

Clarity of Presentation

The demonstration is structured to progressively reveal the system's capabilities:

1. **Problem First**: Start with the pain points travelers face to establish context and relevance
2. **Solution Overview**: Briefly explain the AI-powered approach before the live demo
3. **Show, Don't Tell**: Use the live application to demonstrate every feature rather than just describing it
4. **Explain While Doing**: Narrate each action during the demo to connect UI interactions to underlying technical concepts
5. **Handle Questions**: Pause at key milestones to invite questions

Technical Communication

When explaining technical concepts, use clear analogies and layered explanations:

Technical Concept	Communication Approach
JWT Authentication	"Like a hotel key card — you get it at check-in (login), and every door (API endpoint) checks if it's valid before letting you in"
Row-Level Security	"Like a filing cabinet where each drawer is locked — even database administrators can't accidentally see your data"
LLM Prompt Engineering	"We give the AI a very specific job description with all the rules it must follow — destination, budget, format — and it returns a structured plan"
Pydantic Validation	"Like a bouncer at the door checking IDs — if the AI's JSON output doesn't match our expected format, it gets rejected and we try again"
React Query	"A smart caching layer that remembers recent data so we don't ask the server for the same thing twice"
Fallback Mechanism	"If the primary flight gets cancelled (Gemini fails), there's an automatic rebooking to the next available flight (Groq)"

Demo Communication Guidelines

Do's

- Speak slowly and clearly when showing technical features
- Always explain WHY a feature exists before showing HOW it works
- Acknowledge when AI generation takes time ("The AI is thinking — this is normal at 10-30 seconds")
- Highlight the business value of each feature
- Show error states gracefully if they occur (demonstrates robustness)

Don'ts

- Don't rush through the AI generation wait time — use it to explain the backend process
- Don't skip security/RLS explanation — it's a key differentiator
- Don't show only happy paths — briefly mention what happens on errors (validation messages, retry logic)

Questions Anticipated from Evaluators

Question	Prepared Answer
Why Gemini over GPT-4?	Gemini 2.5 Flash offers structured JSON output natively, fast response times, generous rate limits, and competitive pricing for our use case
How do you handle AI hallucinations?	Pydantic validates every field of the AI output; invalid responses trigger a retry; place names are constrained to the specified destination
What if Supabase goes down?	Supabase provides 99.9% SLA; our health endpoint monitors availability; the application shows graceful error messages
Can multiple users plan the same trip?	Yes — the database has a trip_members table with editor/viewer roles; collaborative features are the next development milestone
How does RLS prevent data leaks?	Every API request uses a Supabase client authenticated with the user's own JWT token; the database's RLS policies filter all queries to only return that user's data
What are the AI generation costs?	Gemini 2.5 Flash is cost-efficient; a typical 3-day itinerary uses ~1000-2000 output tokens; Groq fallback is free tier for development

---

Post-Demo Documentation

All project documentation is available in the GitHub repository:

- https://github.com/Tanay-Ds/Personalized-Travel-Itenerary-Planner

API documentation is always live at:

- https://personalized-travel-itenerary-planner.onrender.com/docs